


.....

Microservices deployment with Kubernetes

.....



Whoami

Containers?

Goals?

Kubernetes?

Capgemini

Decathlon

Cloud Native
Partners

And you?

Agenda



Introduction

Introduction to microservices, and reminder on containerisation.



Kubernetes

Basics on kubernetes, its architecture and these different resources.



Hands-on

Application of the concepts with different labs.

Introduction

Design Microservices Architecture.




“The golden rule: can you make a change
to a service and deploy it by itself
without changing anything else?”

-Sam Newman

Monolith VS Microservices



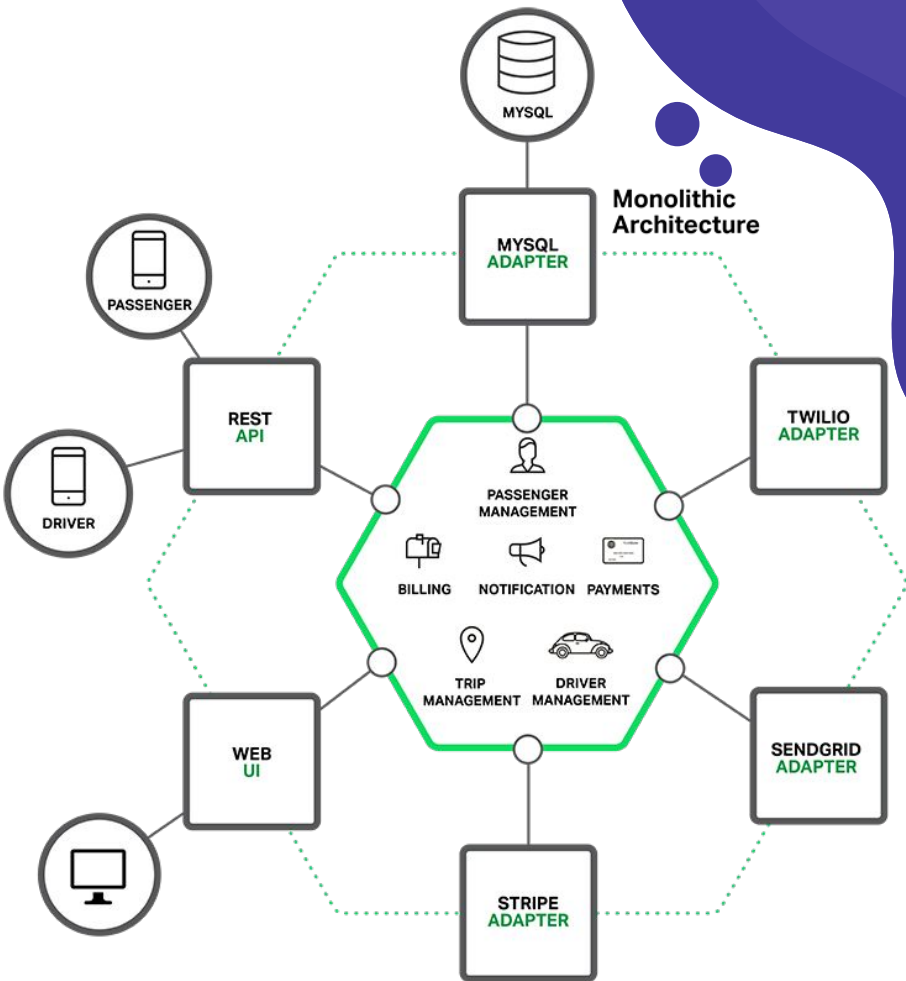
Monolithic
Single Unit



SOA
Coarse-grained



Microservices
Fine-grained



Monolithic application

All the software components of an application are assembled together and tightly packaged.

Challenges with monolithic software

Inflexible

Monolithic applications cannot be built using different technologies.

Unreliable

If even one feature of the system does not work, then the entire system does not work.

Unscalable

Applications cannot be scaled easily since each time the application needs to be updated.

Blocks CD

Many features of an application cannot be built and deployed at the same time.

Slow Development

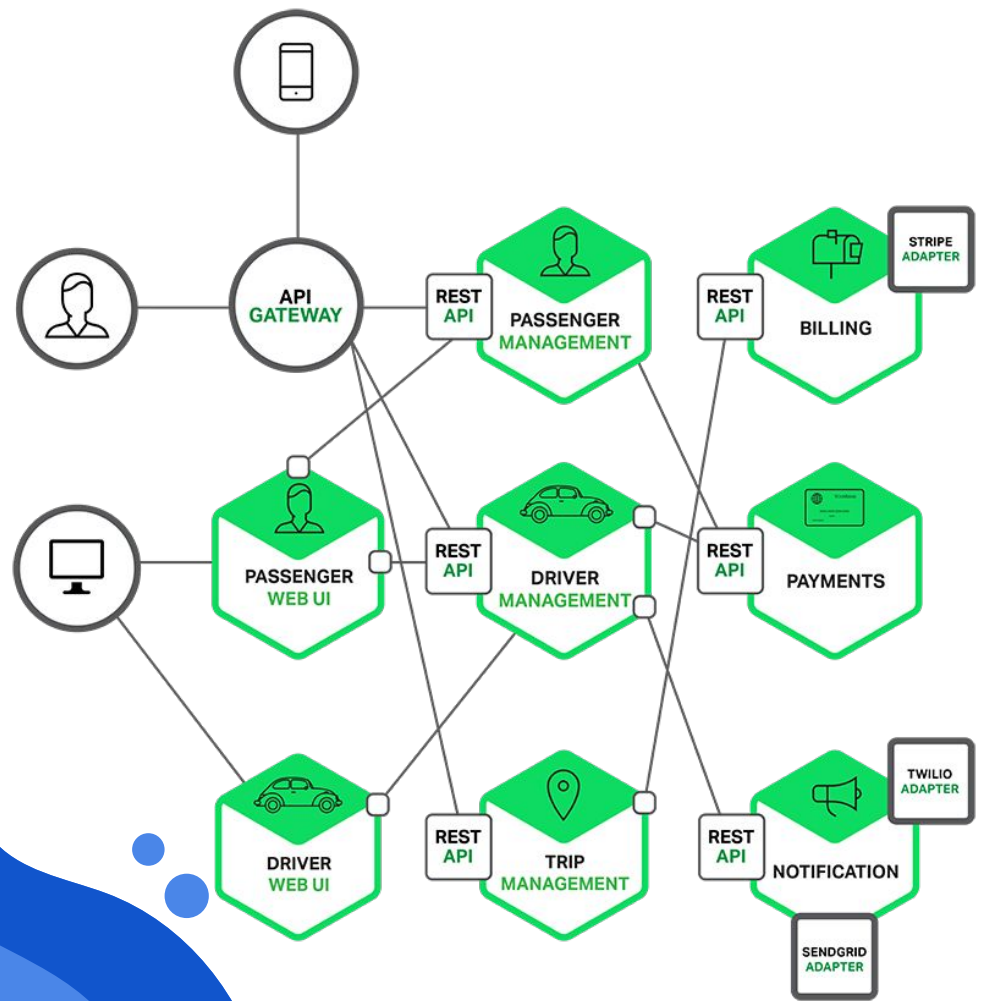
Development in monolithic applications takes a lot of time to be built.

Complexity

Features of complex applications have tightly coupled dependencies.

Microservice application

Each service is self-contained and implements a single business capability.



Features of Microservices.

Decoupling

Services within a system are largely decoupled, so the application can be easily built.

Autonomy

Developers and teams can work independently of each other, thus increasing speed.

Continuous Delivery

Allows frequent releases of software through systematic automation of software creation.

Responsibility

Microservices do not focus on applications as projects.

Decentralized Governance

The focus is on using the right tool for the right job.

Complexity

Microservices support agile development. Any new feature can be quickly developed.

The good, the bad and the...

Pros

- Ability to scale independently
- Fault tolerance
- Can be swapped out or easily re-written
- Framework and language agnostic
- Adheres to KISS principle
- 12 factors compatible

Cons

- More complexity
- Microservice-based architecture may not provide any meaningful benefits
- No greenfield options.
- more robust methods of testing from the entire engineering team
- The need for increased team management and communication.

Summary

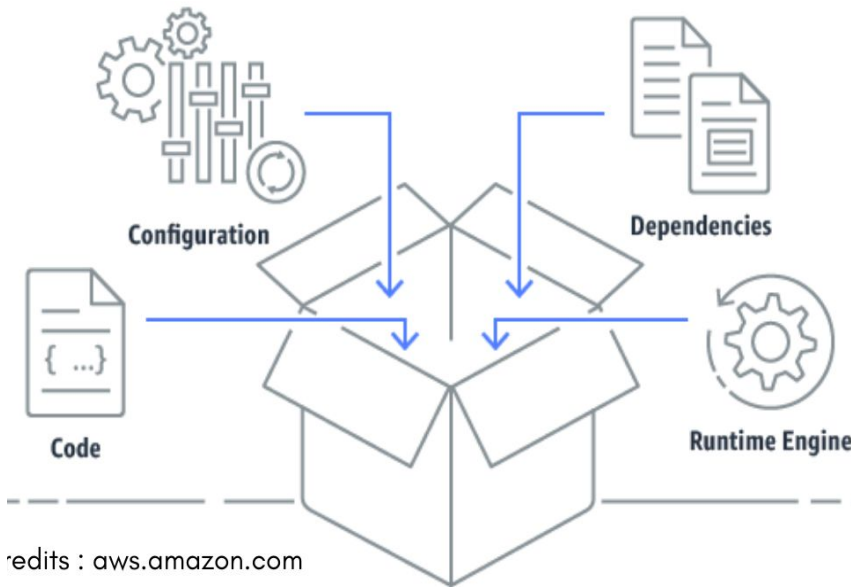
	Pros	Cons
Microservices	Better Organization Decoupled Performance Fewer Mistakes	Cross-cutting Concerns Across Each Service Higher Operational Overhead
Monolith	Fewer Cross-cutting Concerns Less Operational Overhead Performance	Tightly Coupled Harder To Understand

Introduction

Building a microservice in a container

“A container is a software package that contains everything the software needs to run. This includes the executable program as well as system tools, libraries, and settings.”

-Tech Terms




redits : aws.amazon.com

What is a container?

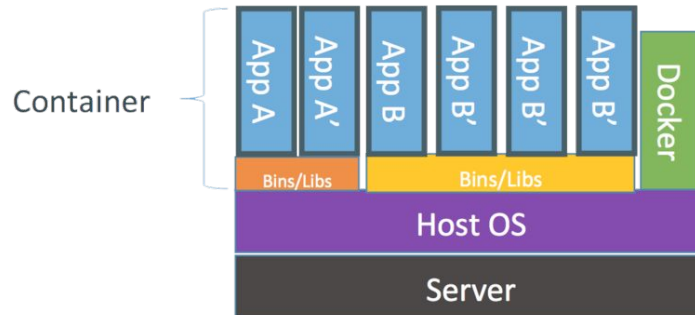
Lightweight, scale-able and isolated VMs in which you run your applications.

Containers VS Virtual machines



Containers are isolated,
but share OS and, where
appropriate, bins/libraries

...result is significantly faster deployment,
much less overhead, easier migration,
faster restart



Why it's important?

Devs

- Build once... (finally) run anywhere
- A clean, safe, hygienic, portable runtime environment for your app.
- No worries about missing dependencies, packages during deployments.
- Automate testing, integration, packaging anything you can script.
- Reduce/eliminate concerns about compatibility on different platforms
- A VM without the overhead of a VM. Instant replay and reset of image snapshots.

Ops

- Configure once... run anything
- Make the entire lifecycle more efficient, consistent, and repeatable
- Eliminate inconsistencies between environments.
- Significantly improves the speed and reliability of continuous deployment and continuous integration systems.
- Better performance, costs, deployment, and portability compare to VMs.

```
FROM ubuntu:16.04
MAINTAINER John Doe <john.doe@example.com>
RUN apt-get update && apt-get install -y python3
RUN python3 --help
```

Dockerfile




WRITE

Define the base image, and describe the different tasks to execute.



build



Docker Image




BUILD

Create your image from your Dockerfile, each step of the script will be executed in order to build the corresponding image.



run

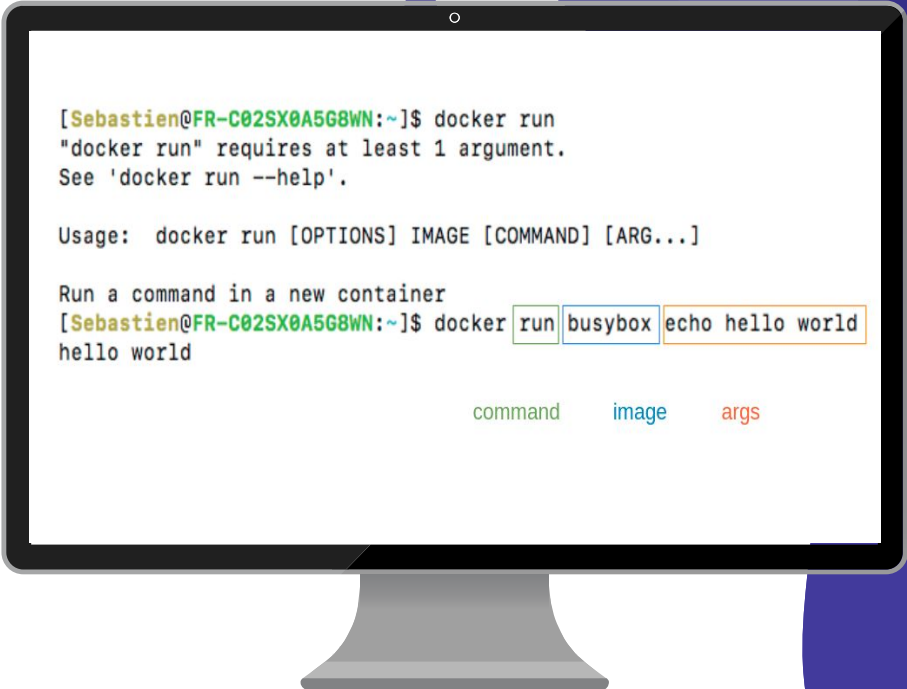


Docker Container



RUN

Start your image on a host, it will execute the main process of the container.



```
[Sebastien@FR-C02SX0A5G8WN:~]$ docker run
"docker run" requires at least 1 argument.
See 'docker run --help'.
```

```
Usage: docker run [OPTIONS] IMAGE [COMMAND] [ARG...]
```

```
Run a command in a new container
```

```
[Sebastien@FR-C02SX0A5G8WN:~]$ docker run busybox echo hello world
hello world
```

command

image

args

Hello World

let's start our first
container!

A computer monitor with a black bezel and a grey stand. The screen displays a terminal window with a white background and black text. The text shows a user running a Docker command to start an Ubuntu container and then checking the OS release information.

```
[Sebastien@FR-C02SX0A5G8WN:~]$ docker run -it ubuntu bash
root@604a05d067fc:/# cat /etc/os-release
NAME="Ubuntu"
VERSION="16.04.3 LTS (Xenial Xerus)"
ID=ubuntu
ID_LIKE=debian
PRETTY_NAME="Ubuntu 16.04.3 LTS"
VERSION_ID="16.04"
HOME_URL="http://www.ubuntu.com/"
SUPPORT_URL="http://help.ubuntu.com/"
BUG_REPORT_URL="http://bugs.launchpad.net/ubuntu/"
VERSION_CODENAME=xenial
UBUNTU_CODENAME=xenial
root@604a05d067fc:/#
```

Let's go inside!

Do not do that in
production

```
[Sebastien@FR-C02SX0A5G8WN:~]$ docker run -d jpetazzo/clock  
b454db0f9bdecb0823f8fb895892075e1580882c039cd90f92e5d2f3deb505f
```

← Run a daemon container

```
[Sebastien@FR-C02SX0A5G8WN:~]$
```

← List active containers

```
[Sebastien@FR-C02SX0A5G8WN:~]$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
b454db0f9bde	jpetazzo/clock	"/bin/sh -c 'while...'"	14 seconds ago	Up 13 seconds

```
[Sebastien@FR-C02SX0A5G8WN:~]$
```

```
[Sebastien@FR-C02SX0A5G8WN:~]$ docker logs b45
```

← Get logs for "b45*"

```
Fri Nov 3 08:42:22 UTC 2017
```

```
Fri Nov 3 08:42:23 UTC 2017
```

```
Fri Nov 3 08:42:24 UTC 2017
```

```
Fri Nov 3 08:42:25 UTC 2017
```

```
[Sebastien@FR-C02SX0A5G8WN:~]$ docker kill b45
```

← Kill container "b45*"

```
b45
```

```
[Sebastien@FR-C02SX0A5G8WN:~]$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
--------------	-------	---------	---------

```
[Sebastien@FR-C02SX0A5G8WN:~]$
```

```
[Sebastien@FR-C02SX0A5G8WN:~]$
```

Let's daemonize!

that's the point of a
container.

[https://github.com/skhedim/epsi-k8s/tree/
master/docker](https://github.com/skhedim/epsi-k8s/tree/master/docker)

Kubernetes

Container Orchestration tool



“Kubernetes is the Linux of the cloud”

-Kelsey Hightower

Load Balancing

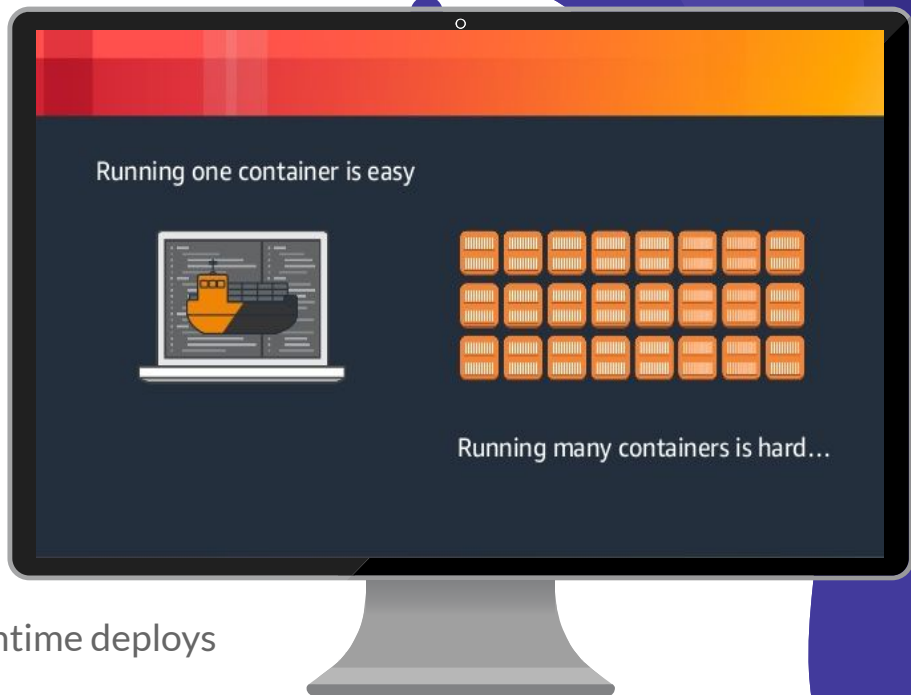
Scaling

Storage

Zero-downtime deploys

Secrets/configuration

High availability



We need to orchestrate

Container orchestration is all about managing the lifecycles of containers, especially in large, dynamic environments.

With an orchestrator you will be able to:



Choose the right place to run containers.

LIFECYCLE



Takes care of restarting the components that were stopped abnormally

FAILOVER



Manage the IPs of the different pods as well as load balancing.

NETWORK



Secrets management
Managing access rights
Configuration files

MUCH MORE



Many tools, but...



Nomad



MESOS



kubernetes

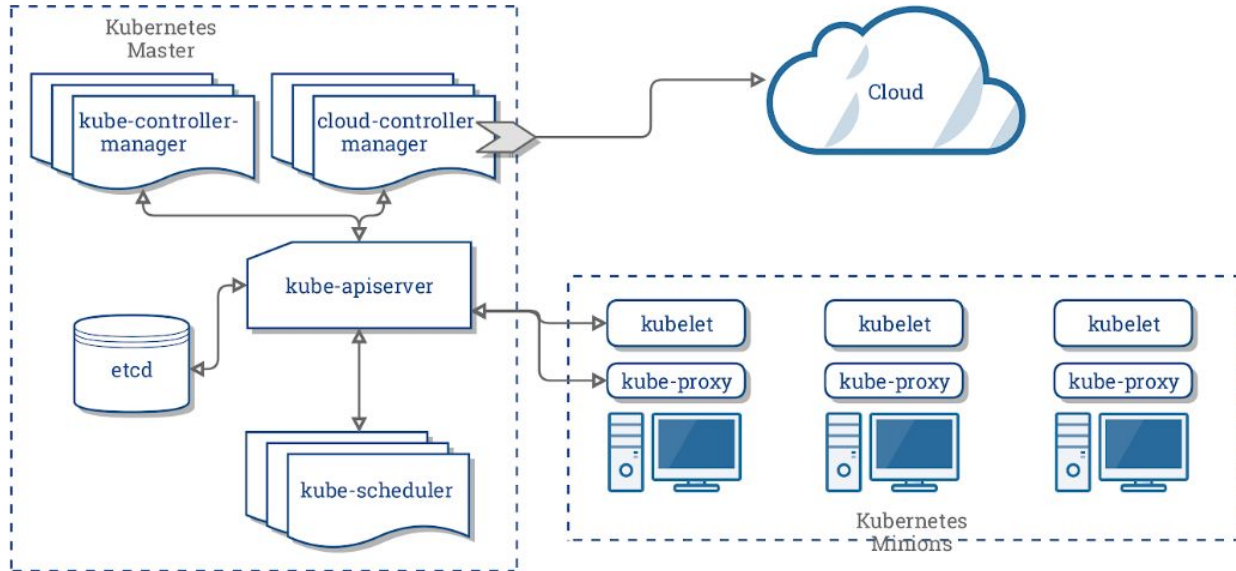
Kubernetes



Architecture

Kubernetes Architecture

Kubernetes itself follows a client-server architecture,

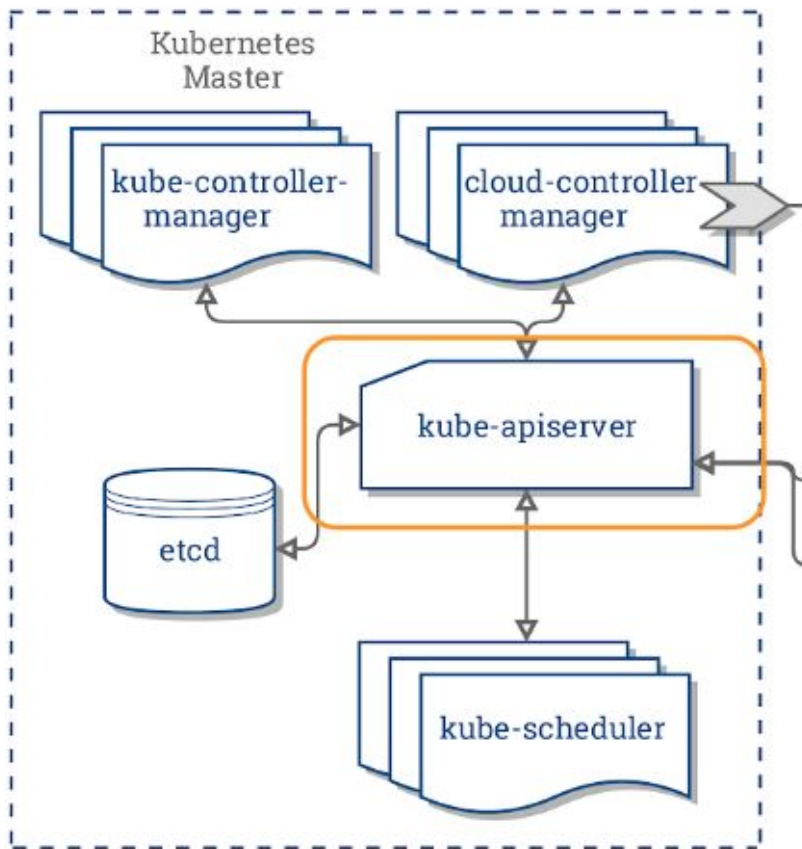


Master:

- ETCD
- kube-apiserver
- kube-controller
- Kube-scheduler
- cloud-controller

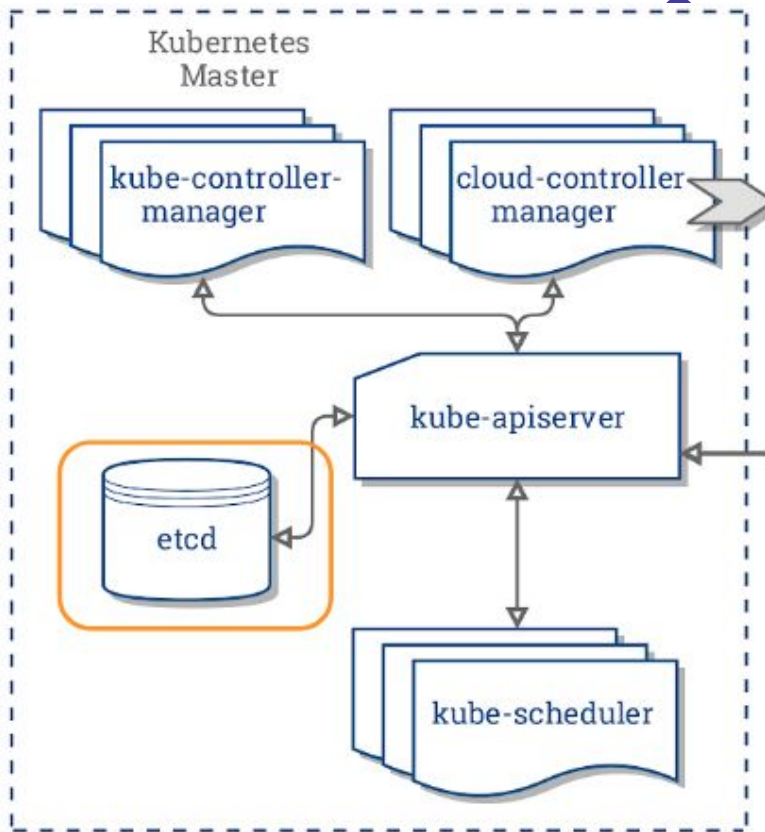
Node:

- kubelet
- kube-proxy



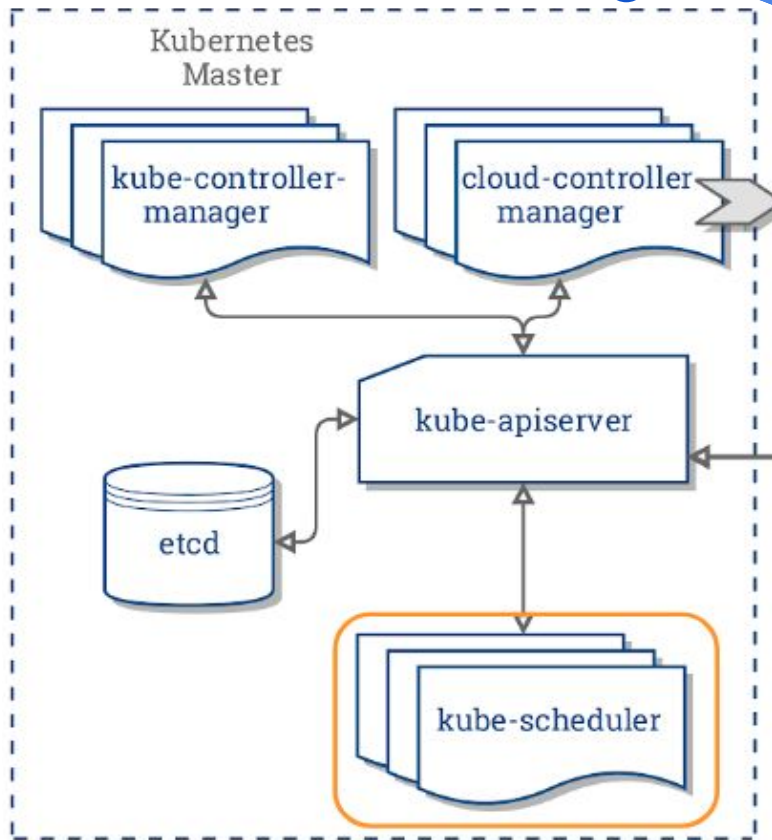
kube-apiserver

Component on the master that exposes the Kubernetes API. It is the front-end for the Kubernetes control plane.



ETCD: Distributed reliable key-value store

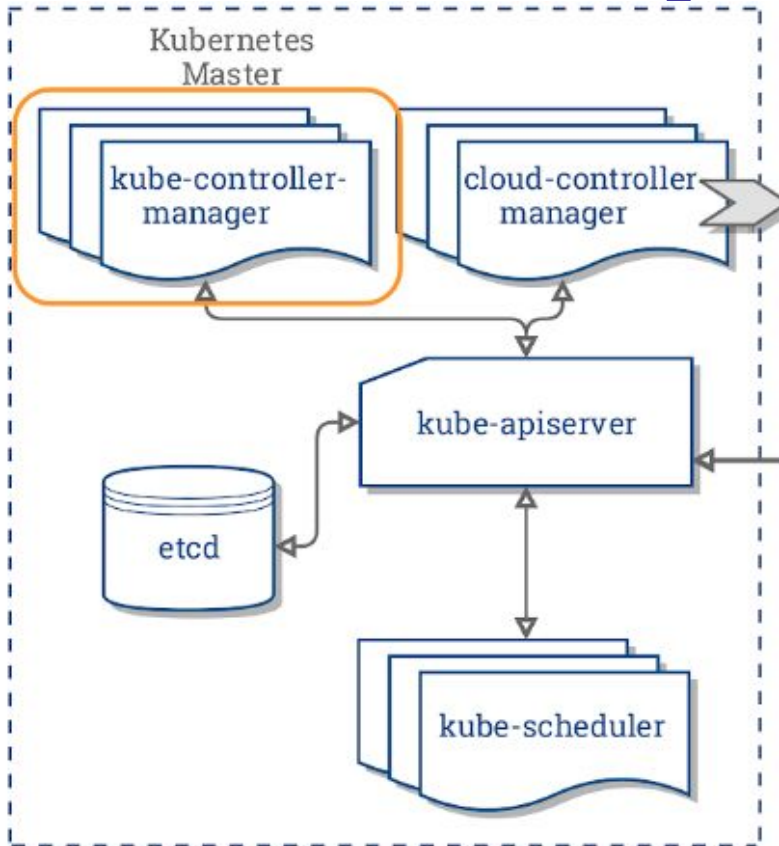
Consistent and highly-available key value store used as Kubernetes' backing store for all cluster data.



kube-scheduler

Component on the master that watches newly created pods that have no node assigned, and selects a node for them to run on.

kube-controllers



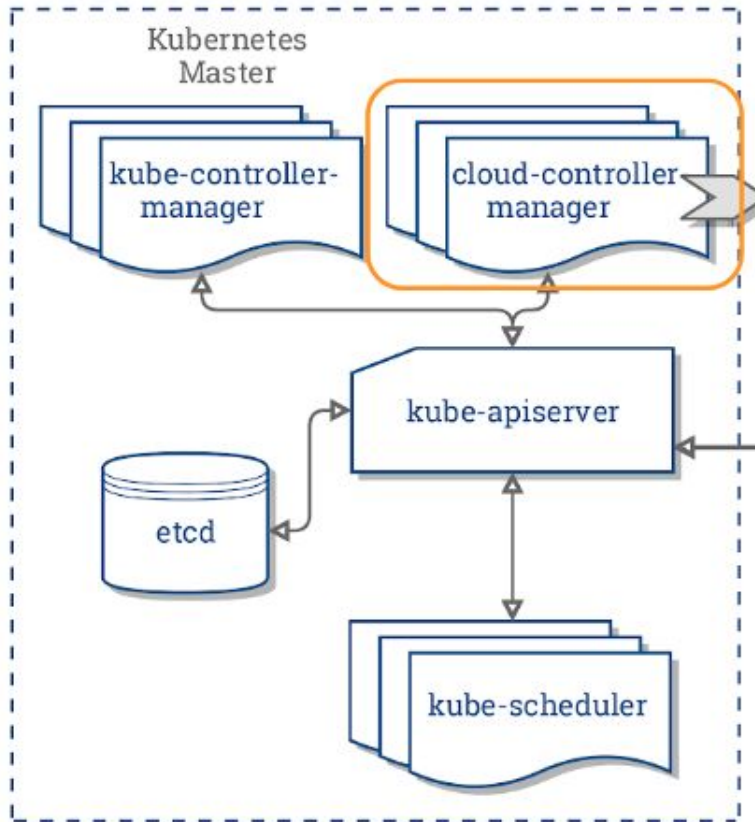
Node Controller: Responsible for noticing and responding when nodes go down.

Replication Controller: Responsible for maintaining the correct number of pods for every replication controller object in the system.

Endpoints Controller: Populates the Endpoints object (that is, joins Services & Pods).

Service Account & Token Controllers: Create default accounts and API access tokens for new namespaces.

cloud-controllers



Node Controller: For checking the cloud provider to determine if a node has been deleted in the cloud after it stops responding

Route Controller: For setting up routes in the underlying cloud infrastructure

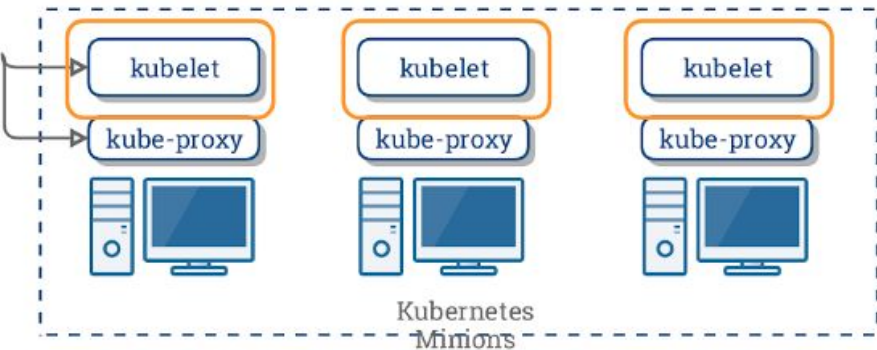
Service Controller: For creating, updating and deleting cloud provider load balancers

Volume Controller: For creating, attaching, and mounting volumes, and interacting with the cloud provider to orchestrate volumes

kubelet


An agent that runs on each node in the cluster. It makes sure that containers are running in a pod.

The kubelet takes a set of PodSpecs that are provided through various mechanisms and ensures that the containers described in those PodSpecs are running and healthy. The kubelet doesn't manage containers which were not created by Kubernetes.



kube-proxy

kube-proxy enables the Kubernetes service abstraction by maintaining network rules on the host and performing connection forwarding.



Kubernetes

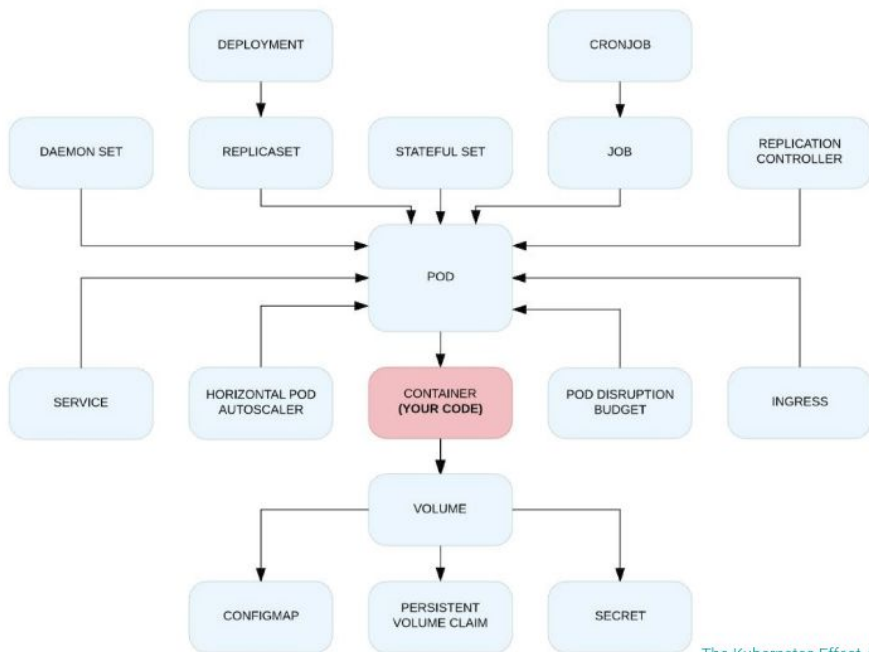
Main features and resources

Kubernetes (commonly stylized as K8s) is an open-source container-orchestration system for automating deployment, scaling and management of containerized applications.

- Wikipedia

Kubernetes objects

Kubernetes objects are persistent entities in the Kubernetes system. Kubernetes uses these entities to represent the state of your cluster.




Pod is the smallest
unit in Kubernetes.

Shared volume

Containerized app

A pod is a group of one or more containers, with shared storage/network, and a specification for how to run the containers.

A pod's contents are always co-located and co-scheduled, and run in a shared context.



YAML File


I hope you like the
YAML.

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod1
labels:
  name: nginx-pod1
spec:
  containers:
  - name: nginx-pod1
    image: nginx:1.7.1
    ports:
    - name: web
      containerPort: 80
```

ReplicaSet is used to maintain a stable set of replica Pods

Used to guarantee the availability of a specified number of identical Pods.

It's recommended to use Deployments instead of directly using ReplicaSets



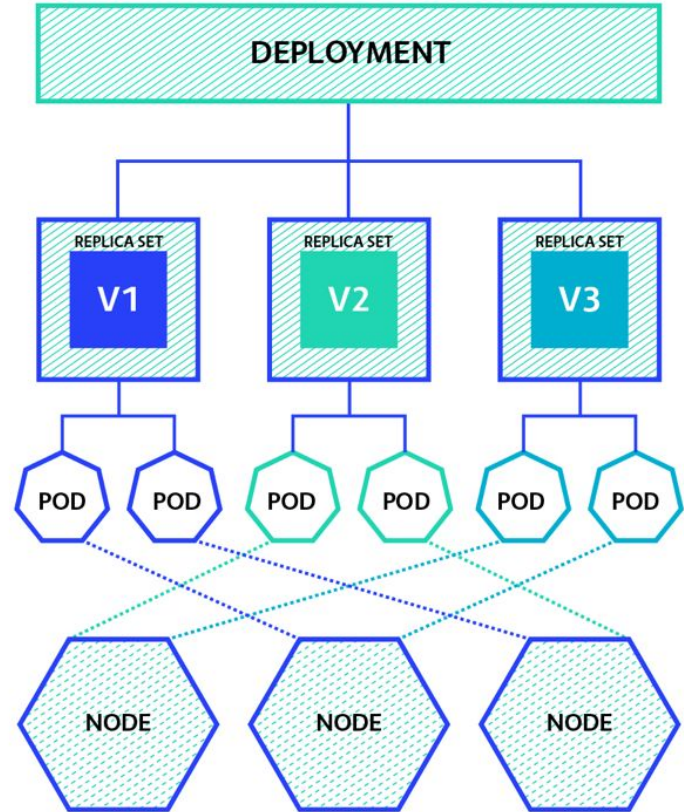
YAML File

Yeah, it's YAML again.

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: frontend
  labels:
    app: guestbook
spec:
  # modify nb of replicas
  replicas: 3
  spec:
    containers:
      - name: nginx
        image : nginx:1.7.1
```

Deployment provides declarative updates for Pods and ReplicaSets.

Service is self-contained and represents a single business capability.




YAML File

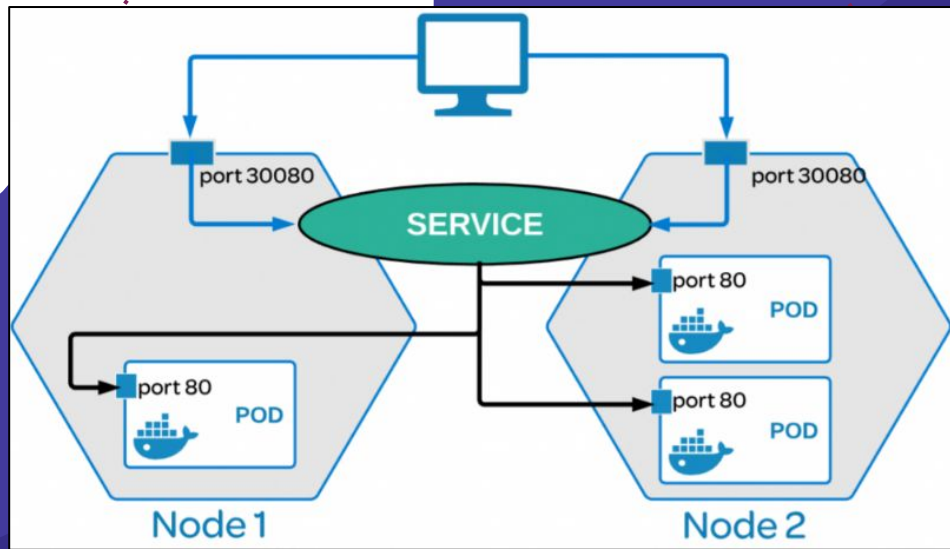
that tablet's getting bigger
and bigger, isn't it?

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```

Network access



An Ingress can be configured to give Services externally-reachable URLs, load balance traffic, terminate SSL / TLS, and offer name based virtual hosting



Persistent Endpoint for Pods.

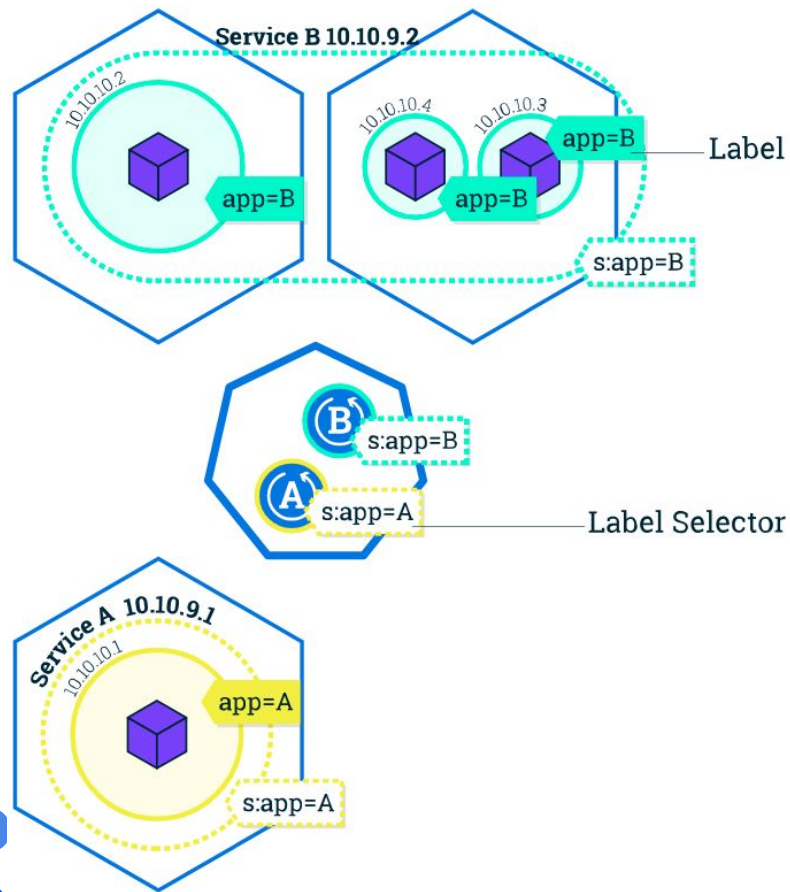
- A Kubernetes Service is an abstraction which defines a logical set of Pods and a policy by which to access them - sometimes called a micro-service.
- Use labels to select Pods

Publishing Services (ServiceTypes)

ClusterIP: Exposes the Service on a cluster-internal IP. Using this value makes the Service only reachable within the cluster.

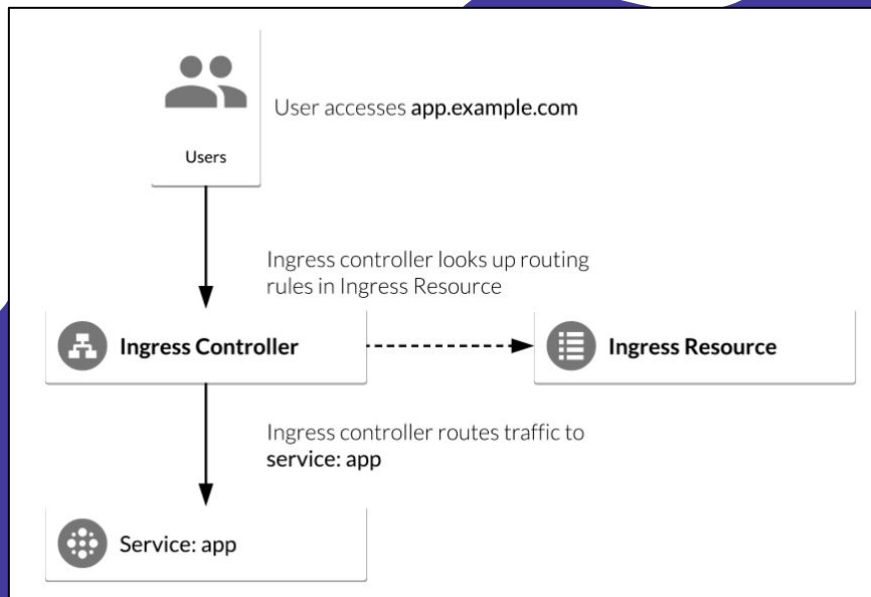
NodePort: Exposes the Service on each Node's IP at a specific port (the NodePort).

LoadBalancer: Exposes the Service externally using a cloud provider's load balancer.



YAML File

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
  clusterIP: 10.0.171.239
  type: LoadBalancer
```



Ingress

An API object that manages external access to the services in a cluster, typically HTTP.

Ingress can provide load balancing, SSL termination and name-based virtual hosting

YAML File

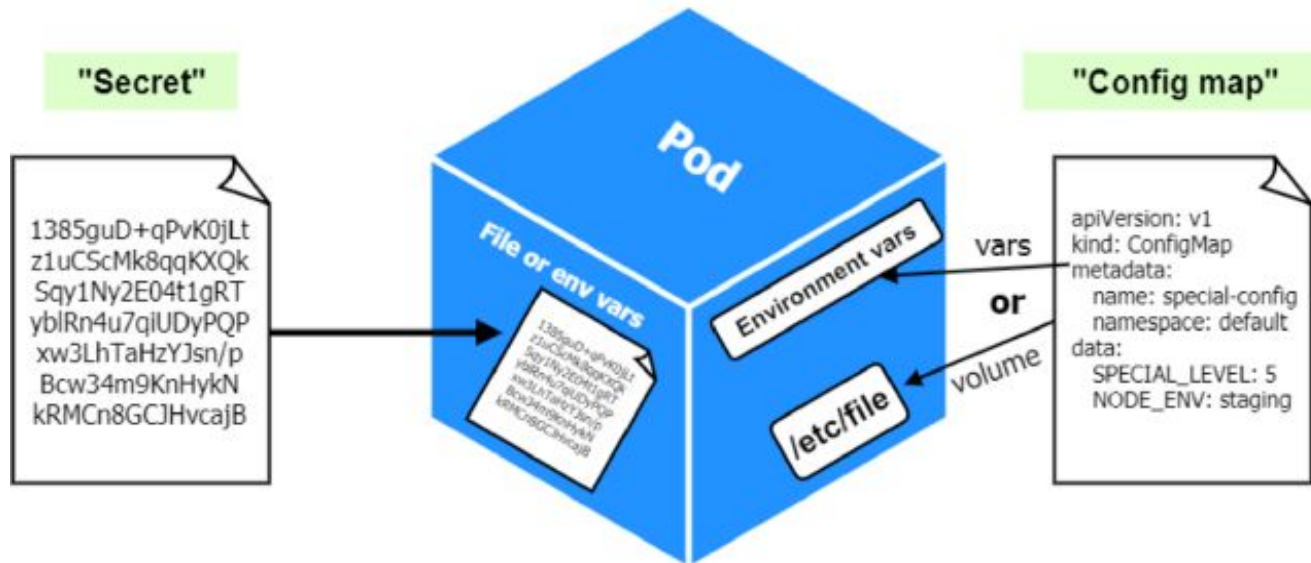
```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: test-ingress
spec:
  rules:
  - http:
      paths:
      - path: /testpath
        backend:
          serviceName: service
          servicePort: 80
```

Secret and Configmaps

Handling Sensitive Information and Container Configurations

Config Maps: a set of values that can be mapped to a pod as “volume” or passed as environment variables.

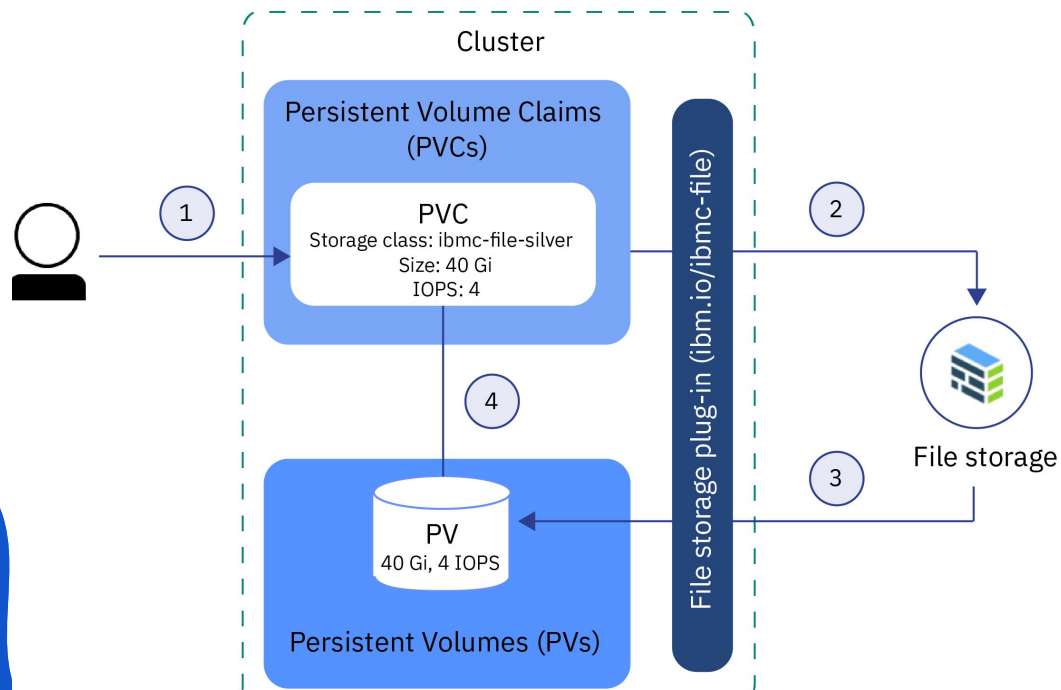
Secrets: similar to config maps, secrets can be mounted into a pod as a volume to expose needed information or can be injected as environment variables.



Persist data

PersistentVolume (PV): resources in the cluster. A storage in the cluster that has been provisioned by an administrator or dynamically provisioned using Storage Classes.

PersistentVolumeClaim (PVC): Request for storage by a user. It is similar to a pod. Pods consume node resources and PVCs consume PV resources.



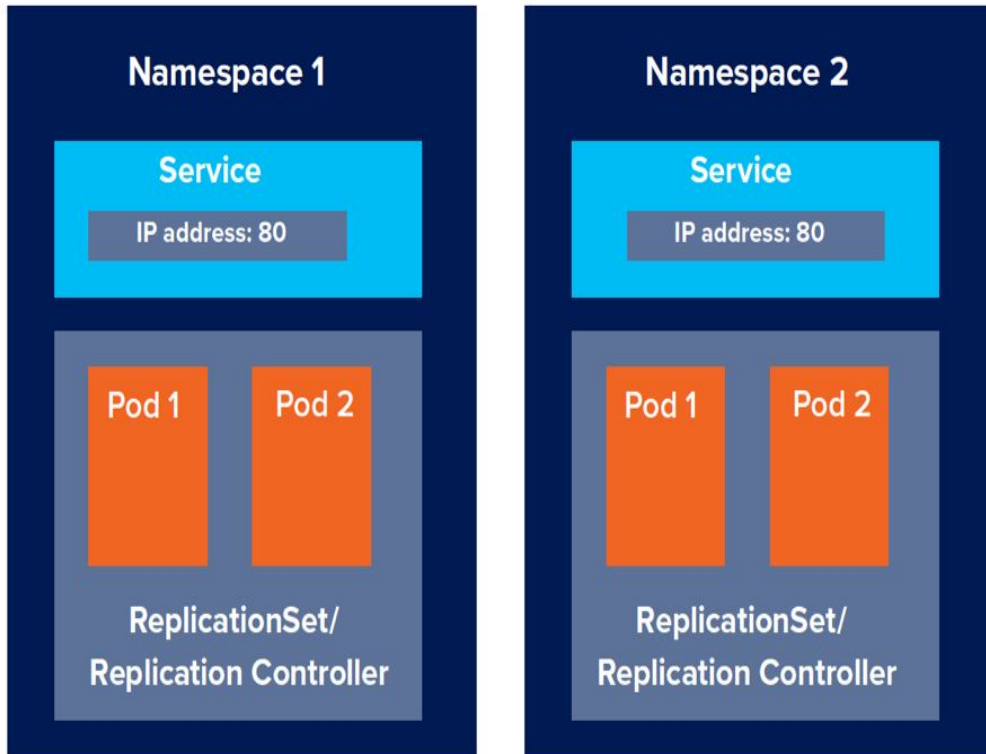
PV

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: task-pv-volume
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/mnt/data"
```

PVC

```
apiVersion: v1
kind:
  PersistentVolumeClaim
metadata:
  name: myclaim
spec:
  accessModes:
    - ReadWriteOnce
  volumeMode: Filesystem
  resources:
    requests:
      storage: 8Gi
  storageClassName:
    my-storage-class
```

Namespaces are a logical isolation of the different resources seen previously



Use-case #1: Roles and Responsibilities in an Enterprise

Use-case #2: Using Namespaces to partition development landscapes

Use-case #3: Partitioning of your Customers



So many more!

ClusterRole

ServiceAccounts

mutatingwebhookconfigurations

StorageClass

CronJobs

DaemonSets

PodDisruptionBudget

ClusterRoleBinding

CRDs

Network policies

Statefulsets



kubectl



Kubernetes

K8s API

But wait! How can I
deploy all of these
objects?

kubectl, a CLI tool to manage your cluster.

- works with a config file
- talks with the API server
- only one tool to do everything

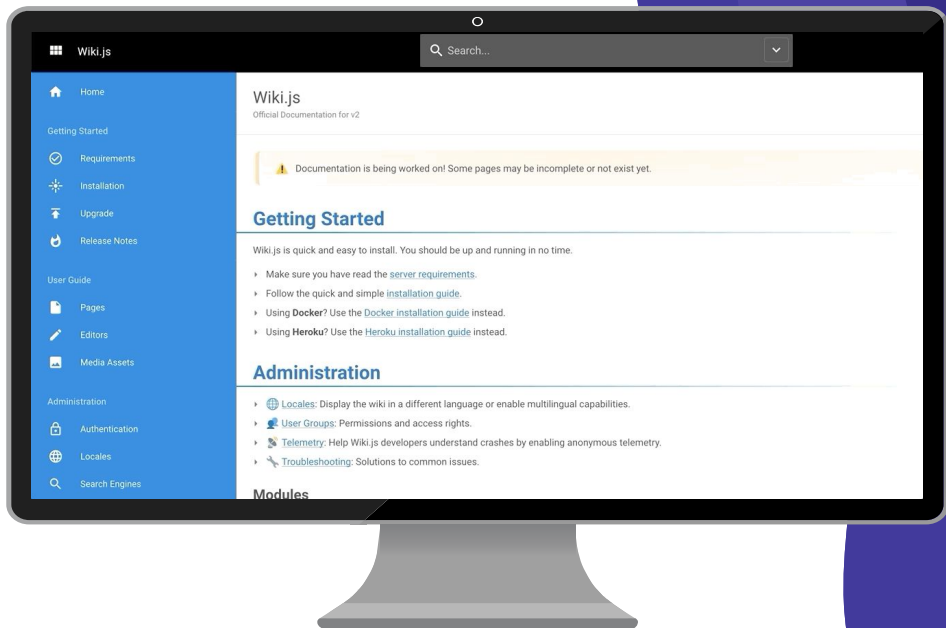


[https://github.com/skhedim/epsi-k8s/tree/master/
kubernetes](https://github.com/skhedim/epsi-k8s/tree/master/kubernetes)



Hands-on

Package an application

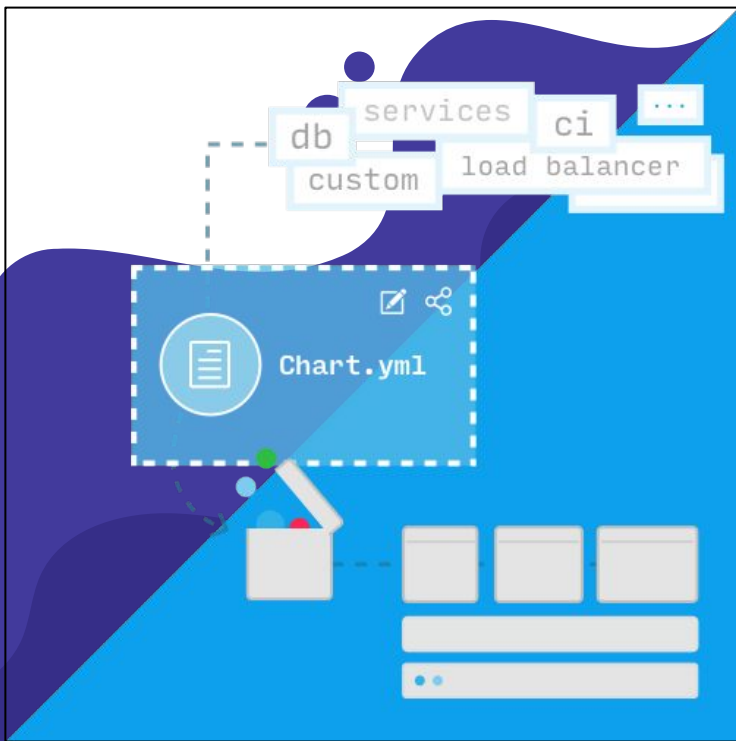


Wiki.js

A wiki engine running on Node.js and written in JavaScript.

There is a docker image of the application

But nothing is packaged for Kubernetes



Package your objects

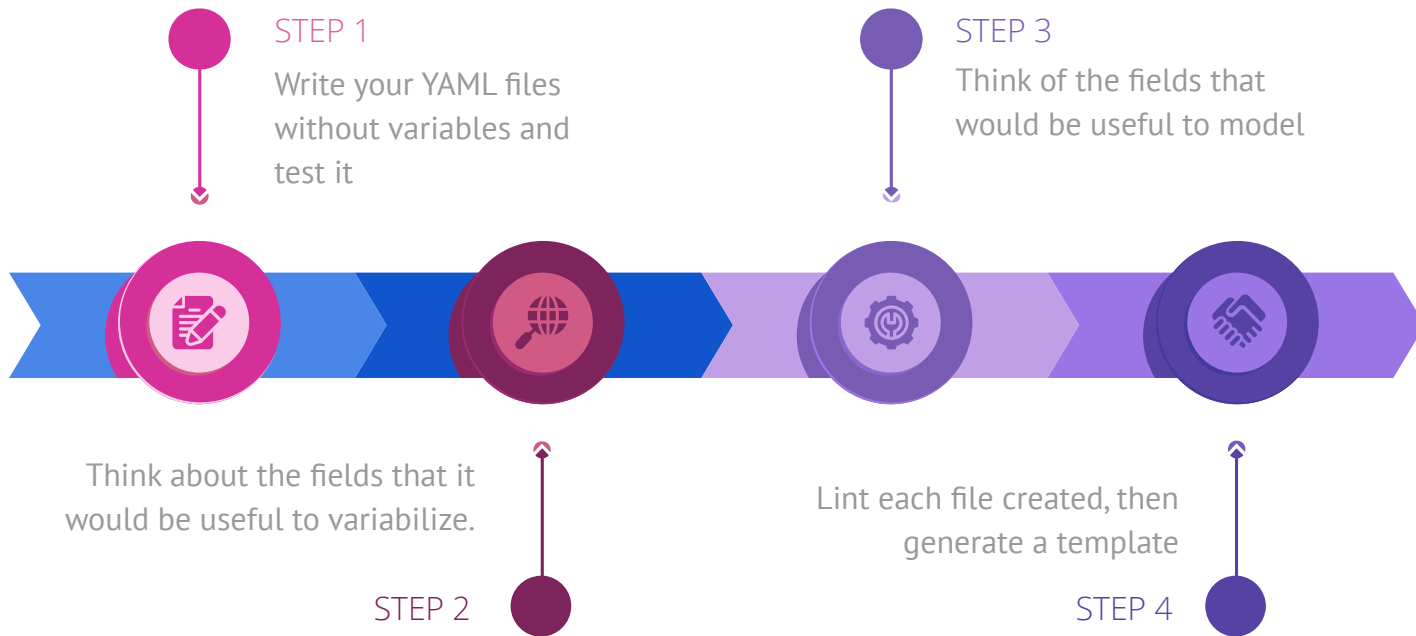
Helm packages multiple Kubernetes resources into a single object

logical deployment called chart

A Chart is package, a set of k8s resources (deployments, configmaps, services, ingress,...)

A Template is a Kubernetes resource model based on the Golang template engine and the Sprig library

Think about it!



[https://github.com/skhedim/epsi-k8s/tree/
master/helm](https://github.com/skhedim/epsi-k8s/tree/master/helm)



Thanks!

Does anyone have any questions?